

Long Naïve Bayes question

Question

You have a set of book reviews. In the left column of the following table, each of "a", "b", "c" represents a different word. In the right column, "1" means the review is positive and "0" otherwise.

Review Text	Positive
aaab	1
cbaba	1
aba	1
cb	1
bccb	0
acc	0
abbc	0

If you want to use the review text to predict its tone (positive or not) with Naive Bayes Classification:

i) What assumption does Naive Bayes make on the data?

Assuming the naive assumption is met, and you have an unlabeled piece of review text to classify: "bc". When answering the following questions, please write down the steps.

ii) Compute the prior $p(\text{tone}=1)$, $p(\text{tone}=0)$, $p(b)$, and $p(c)$. DO NOT use any smoothing method. Present the answer as fractions.

iii) Compute the posterior $p(\text{tone}=1|bc)$ and $p(\text{tone}=0|bc)$. Present the answer as decimals with 3 digit precision. What is the tone of this review?

Answer

- Feature attributes are conditionally independent conditioned on the label; (0.5 pts)
- priors:
 - $P(\text{tone}=1) = 4/7$;
 - $P(\text{tone}=0) = 3/7$
 - $P(b) = 9/25$
 - $P(c) = 7/25$
 - (1.5 pts)
- Posteriors:

$$P(bc) = P(\text{tone} = 1) P(b|\text{tone} = 1) P(c|\text{tone} = 1) + P(\text{tone} = 0) P(b|\text{tone} = 0) P(c|\text{tone} = 0)$$

$$= \frac{4}{7} \times \frac{5}{14} \times \frac{2}{14} + \frac{3}{7} \times \frac{4}{11} \times \frac{5}{11}$$

$$P(\text{tone} = 1 | bc) = \frac{P(\text{tone}=1) P(b|\text{tone} = 1) P(c|\text{tone} = 1)}{P(bc)} = \frac{\frac{4}{7} \times \frac{5}{14} \times \frac{2}{14}}{\frac{4}{7} \times \frac{5}{14} \times \frac{2}{14} + \frac{3}{7} \times \frac{4}{11} \times \frac{5}{11}} = 0.292 \quad (0.4$$

pts; OK if differ slightly in the last digits)

$$P(\text{tone} = 0 | bc) = \frac{P(\text{tone}=0) P(b|\text{tone} = 0) P(c|\text{tone} = 0)}{P(bc)} = \frac{\frac{3}{7} \times \frac{4}{11} \times \frac{5}{11}}{\frac{4}{7} \times \frac{5}{14} \times \frac{2}{14} + \frac{3}{7} \times \frac{4}{11} \times \frac{5}{11}} = 0.708 \quad (0.4 \text{ pts;}$$

OK if differ slightly in the last digits)

negative because $P(\text{tone} = 0 | bc) > P(\text{tone} = 1 | bc)$ (0.2 pts)

=====

Question

You have a set of book reviews. In the left column of the following table, each of "a", "b", "c" represents a different word. In the right column, "1" means the review is positive and "0" otherwise.

Review Text	Positive
baa	0
bcaba	0
abaa	0
bc	0
babc	1
cac	1
bbcc	1

If you want to use the review text to predict its tone (positive or not) with Naive Bayes Classification:

i) What assumption does Naive Bayes make on the data?

Assuming the naive assumption is met, and you have an unlabeled piece of review text to classify: "cb". When answering the following questions, please write down the steps. If the answers are fractions, you can but don't need to convert them to decimals.

ii) Compute the prior $p(\text{tone}=1)$, $p(\text{tone}=0)$, $p(b)$, and $p(c)$. DO NOT use any smoothing method. Present the answer as fractions.

iii) Compute the posterior $p(\text{tone}=1|cb)$ and $p(\text{tone}=0|cb)$. Present the answer as decimals with 3 digit precision. What is the tone of this review?

Answer:

- Feature attributes are conditionally independent conditioned on the label; (0.5 pts)

- priors:

$$P(\text{tone}=1) = 3/7;$$

$$P(\text{tone}=0) = 4/7$$

$$P(b) = 9/25$$

$$P(c) = 7/25$$

(1.5 pts)

- Posteriors:

$$P(cb) = P(\text{tone} = 1) P(b|\text{tone} = 1) P(c|\text{tone} = 1)$$

$$+ P(\text{tone} = 0) P(b|\text{tone} = 0) P(c|\text{tone} = 0) = \frac{3}{7} \times \frac{4}{11} \times \frac{5}{11} + \frac{4}{7} \times \frac{5}{14} \times \frac{2}{14}$$

$$P(\text{tone} = 1 | cb) = \frac{P(\text{tone}=1) P(b|\text{tone} = 1) P(c|\text{tone} = 1)}{P(cb)} = \frac{\frac{3}{7} \times \frac{4}{11} \times \frac{5}{11}}{\frac{3}{7} \times \frac{4}{11} \times \frac{5}{11} + \frac{4}{7} \times \frac{5}{14} \times \frac{2}{14}} = 0.708 \quad (0.4$$

pts; OK if differ slightly in the last digits)

$$P(\text{tone} = 0 | cb) = \frac{P(\text{tone}=0) P(b|\text{tone} = 0) P(c|\text{tone} = 0)}{P(cb)} = \frac{\frac{4}{7} \times \frac{5}{14} \times \frac{2}{14}}{\frac{3}{7} \times \frac{4}{11} \times \frac{5}{11} + \frac{4}{7} \times \frac{5}{14} \times \frac{2}{14}} = 0.292 \quad (0.4$$

pts; OK if differ slightly in the last digits)

positive because $P(\text{tone} = 1 | bc) > P(\text{tone} = 0 | bc)$ (0.2 pts)

Long CNN:

Q:

Your input is a 58 * 91 RGB image, and you use a convolutional layer with 100 filters that are each 6 * 5 and with one bias parameter for each filter, using a stride of (4,6) and a total padding of 4.

1. What is the output shape?

2. What is the number of parameters (including the bias parameters)?

3. Suppose we further use a 3 * 4 max-pooling layer (with a stride of (3,4) and no padding) after the above convolutional layer. What is the output shape of the max-pooling layer?

Answer: (1) $15 * 16 * 100$
 (2) 9100
 (3) $5 * 4 * 100$

Alternative Q:

Your input is a $131 * 63$ RGB image, and you use a convolutional layer with 100 filters that are each $7 * 4$ and with one bias parameter for each filter, using a stride of (4,7) and a total padding of 4.

1. What is the output shape?
2. What is the number of parameters (including the bias parameters)?
3. Suppose we further use a $3 * 2$ max-pooling layer (with a stride of (3,2) and no padding) after the above convolutional layer. What is the output shape of the max-pooling layer?

Answer: (1) $33 * 10 * 100$
 (2) 8500
 (3) $11 * 5 * 100$

Long RL question:

Version 1

Q1 (1.5pts)

$V^*(S1) = R(S1)=10$, $V^*(S2) = R(S2)=20$. (0.5pts)

Expected utility for different actions:

Action A1: $V = 5 + 0.8(0.1 * 10 + 0.9 * 20) = 20.2$

Action A2: $V = 5 + 0.8(0.9 * 10 + 0.1 * 20) = 13.8$

Action A3: $V = 5 + 0.8(0.5 * 10 + 0.5 * 20) = 17$

Therefore, the optimal policy from state S0 is to take Action A1.
(expected utility + optimal action: 1pts)

Q2 (0.5pts)

The notable difference between Q-learning and SARSA is that when updating $Q(s,a)$: Q-learning uses the expected cumulative discounted reward for taking an **optimal** action from state s' whereas SARSA uses the expected cumulative discounted reward for the **actual** action it took from state s' . Also, SARSA also updates $Q(s', a')$ in each iteration.

Q3(0.5pts)

The Epsilon-greedy algorithm enables a balance between exploration and exploitation, so that both SARSA and Q-learning could bootstrap with non-initial value (stuck).

Q4 (0.5pts)

The max is 8: each state leads to an update to a distinct entry.

The min is 1: all observations lead to updates to the same (s,a) entry.

Version 2

Q1(1.5pts)

$V^*(S1) = R(S1)=10$, $V^*(S2) = R(S2)=25$. (0.5pts)

Expected utility for different actions:

$$\text{Action A1: } 10 + 0.9(0.8 * 10 + 0.2 * 25) = 21.7$$

$$\text{Action A2: } V = 10 + 0.9(0.2 * 10 + 0.8 * 25) = 29.8$$

$$\text{Action A3: } V = 10 + 0.9(0.5 * 10 + 0.5 * 25) = 25.75$$

Therefore, the optimal policy from state S0 is to take Action A2.
(expected utility + optimal action: 1pts)

Q2 (0.5pts)

The notable difference between Q-learning and SARSA is that when updating $Q(s,a)$: Q-learning uses the expected cumulative discounted reward for taking an **optimal** action from state s' whereas SARSA uses the expected cumulative discounted reward for the **actual** action it took from state s' . Also, SARSA also updates $Q(s', a')$ in each iteration.

Q3 (0.5pts)

The Epsilon-greedy algorithm enables a balance between exploration and exploitation, so that both SARSA and Q-learning could bootstrap with non-initial value (stuck).

Q4 (0.5pts)

The max is 8: each state leads to an update to a distinct entry.

The min is 1: all observations lead to updates to the same (s,a) entry.

Long Search question:

Version 1: Double star graph:

A BFS will start at 0 and then consider every node one edge away, which are 1,2,...,n. By the tie-breaking rule, it will follow in that order. Then, after node n is visited, every node in the second layer is visited, which are n+1, n+2, ..., 2n.

A DFS will start at 0 then visit 1 then 2 at which point it has to back up back to 0. This process is repeated: visiting i then i+n from 0 and then continuing on with i+1 till all nodes are visited, eventually ending with n and 2n.

Version 2: wheel graph:

A BFS will start at 0 and then consider every node one edge away, which are 1, 2..., n. By the tie-breaking rule, it will follow in that order.

A DFS will start at 0 then visit 1 at which point it follows the outer "cycle" 1,2,3...,n.